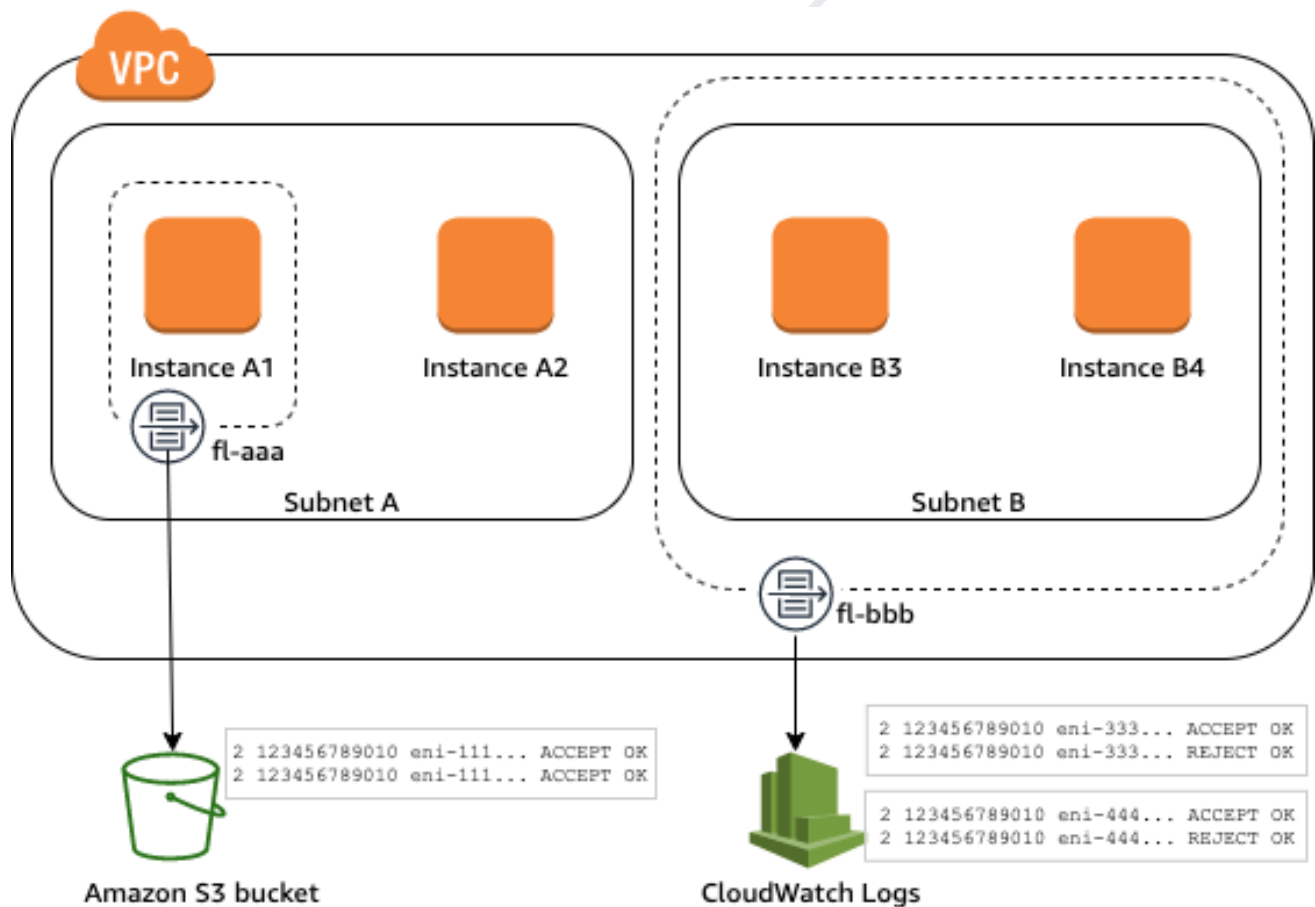


## #12 Task: Create VPC Flow Logs, S3 Access Logs, API Gw Logs, Lambda Logs, CloudTrail logs and Install CloudWatch Monitoring Agents on EC2 to capture data for logs. All these Logs can be used further to configure a Cloud Vulnerability Scanning System to maintain a robust & secured Infra Environment.

### 1. Setup VPC Flow Logs To CloudWatch Log Group Or S3 Bucket.



## 1. Go to your VPC Dashboard on the AWS console

You are using the following Amazon VPC resources

VPCs  
See all regions ▼ Ohio 1

NAT Gateways  
See all regions ▼ Ohio 0

Subnets  
See all regions ▼ Ohio 3

VPC Peering Connections  
See all regions ▼ Ohio 0

Route Tables  
See all regions ▼ Ohio 1

Network ACLs  
See all regions ▼ Ohio 1

Internet Gateways  
See all regions ▼ Ohio 1

Security Groups  
See all regions ▼ Ohio 5

Egress-only Internet Gateways  
See all regions ▼ Ohio 0

Customer Gateways  
See all regions ▼ Ohio 0

[View complete service health details](#)

### Account Attributes

[Resource ID length management](#)

### Additional Information

[VPC Documentation](#)  
[All VPC Resources](#)  
[Forums](#)  
[Report an Issue](#)

### Transit Gateway Network Manager

Network Manager enables centrally manage your global network across AWS and on-premises. [Learn more](#)

## 2. Click on Flow logs

Description

CIDR Blocks

Flow Logs

Tags

You can create flow logs on your resources to capture IP traffic flow information for the network interfaces for your resources. [Learn more](#)

Create flow log

Actions ▼

None found

Flow Log ID

Filter

Destination Ty

Destination Name

IAM Role ARN

You do not have any Flow Logs in this region

## Select Flow logs in the VPC Dashboard

## 3. Click on Create Flow logs

### Create flow log

Flow logs can capture IP traffic flow information for the network interfaces associated with your resources. You can create multiple subscriptions to send traffic to different destinations. [Learn more](#)

Resources vpc-b9d638d2 ⓘ

Filter All ⓘ

Maximum aggregation interval  
☒ 10 minutes ⓘ  
☐ 1 minute

Destination  
☒ Send to CloudWatch Logs ⓘ  
☐ Send to an S3 bucket

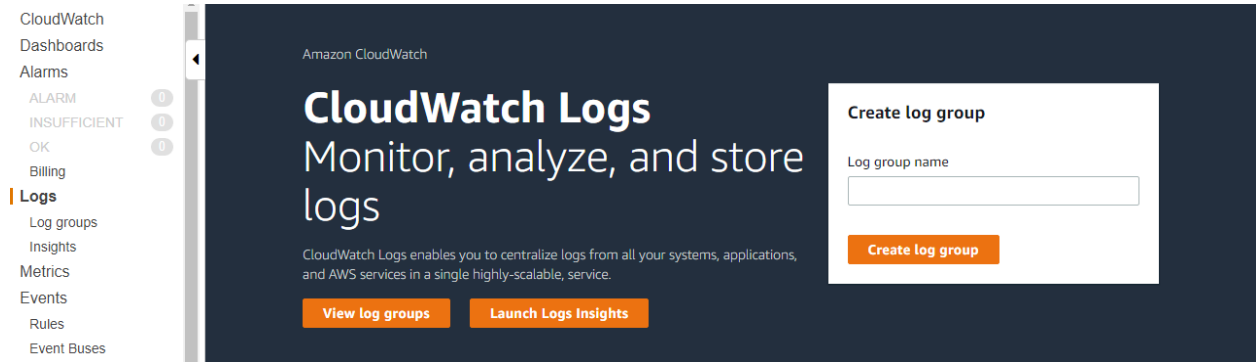
Destination log group\* Enter log destination ⓘ

IAM role\* No IAM role selected ⓘ

## Create Flow log

If you don't have any Destination log group or IAM Role create one from scratch.

- To Create a Destination log group, Go to the Cloudwatch dashboard on your VPC Console

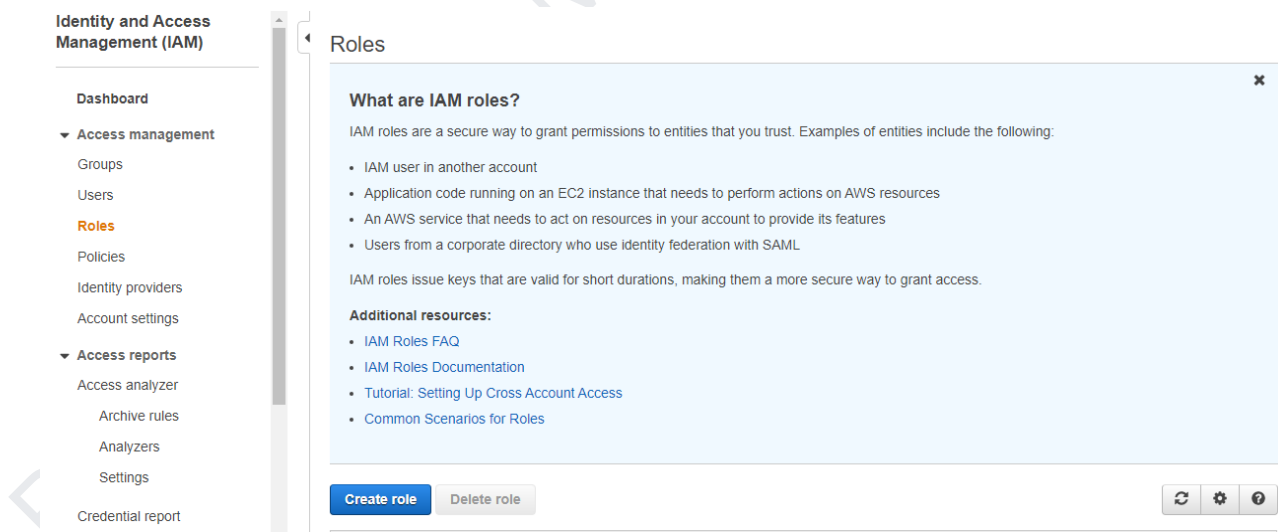


Click on Logs and enter a log group name, click on create log

4. Return back to the VPC Flowlogs dashboard and enter the Log group name created.

5. Create an IAM role which will allow VPC to Write to the log group

- To create an IAM role, go to the Identity and Access Management dashboard.



Click on Create a Role.

- The role that will be created will be an EC2 Role.

**AWS service**  
EC2, Lambda and others

**Another AWS account**  
Belonging to you or 3rd party

**Web identity**  
Cognito or any OpenID provider

**SAML 2.0 federation**  
Your corporate directory

Allows AWS services to perform actions on your behalf. [Learn more](#)

### Choose a use case

#### Common use cases

##### EC2

Allows EC2 instances to call AWS services on your behalf.

##### Lambda

Allows Lambda functions to call AWS services on your behalf.

#### Or select a service to view its use cases

<a href="#">API Gateway</a>	<a href="#">CodeGuru</a>	<a href="#">ElastiCache</a>	<a href="#">Kinesis</a>	<a href="#">RoboMaker</a>
<a href="#">AWS Backup</a>	<a href="#">CodeStar Notifications</a>	<a href="#">Elastic Beanstalk</a>	<a href="#">Lake Formation</a>	<a href="#">S3</a>
<a href="#">AWS Chatbot</a>	<a href="#">Comprehend</a>	<a href="#">Elastic Container Service</a>	<a href="#">Lambda</a>	<a href="#">SMS</a>
<a href="#">AWS Support</a>	<a href="#">Config</a>	<a href="#">Elastic Transcoder</a>	<a href="#">Lex</a>	<a href="#">SNS</a>
<a href="#">Amplify</a>	<a href="#">Connect</a>	<a href="#">ElasticLoadBalancing</a>	<a href="#">License Manager</a>	<a href="#">SWF</a>

\* Required

Cancel

Next: Permissions

### Select EC2 Role.

- Click Next to enter Tag names

### Create role

1 2 3 4

#### ▼ Attach permissions policies

Choose one or more policies to attach to your new role.

Create policy

Filter policies

Search

Showing 664 results

	Policy name	Used as
<input type="checkbox"/>	<a href="#">AlexaForBusinessDeviceSetup</a>	None
<input type="checkbox"/>	<a href="#">AlexaForBusinessFullAccess</a>	None
<input type="checkbox"/>	<a href="#">AlexaForBusinessGatewayExecution</a>	None
<input type="checkbox"/>	<a href="#">AlexaForBusinessNetworkProfileServicePolicy</a>	None
<input type="checkbox"/>	<a href="#">AlexaForBusinessPolyDelegatedAccessPolicy</a>	None
<input type="checkbox"/>	<a href="#">AlexaForBusinessReadOnlyAccess</a>	None
<input type="checkbox"/>	<a href="#">AmazonAPIGatewayAdministrator</a>	None

\* Required

Cancel

Previous

Next: Tags

### Click on Next:Tags

- Enter the Role name and click on create a role

## Create role

1 2 3 4

### Review

Provide the required information below and review this role before you create it.

Role name\* vpc-flow-logs-demo-role

Use alphanumeric and '+=, @-\_' characters. Maximum 64 characters.

Role description Allows EC2 instances to call AWS services on your behalf.

Maximum 1000 characters. Use alphanumeric and '+=, @-\_' characters.

Trusted entities AWS service: ec2.amazonaws.com

Policies Policies not attached

Permissions boundary Permissions boundary is not set

No tags were added.

\* Required

Cancel

Previous

Create role

- Attach the IAM roles for publishing flow logs to CloudWatch Logs
- Go to the Role you created and attach the inline policy

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Action": [
        "logs:CreateLogGroup",
        "logs:CreateLogStream",
        "logs:PutLogEvents",
        "logs:DescribeLogGroups",
        "logs:DescribeLogStreams"
      ],
      "Effect": "Allow",
      "Resource": "*"
    }
  ]
}
```

- Ensure that your role has a trust relationship that allows the flow logs service to assume the role. This allows EC2 to write into the Log-group.
- Click on Trust Relationship in the Roles Dashboard

Permissions Trust relationships Tags Access Advisor Revoke sessions

You can view the trusted entities that can assume the role and the access conditions for the role. [Show policy document](#)

[Edit trust relationship](#)

**Trusted entities**

The following trusted entities can assume this role.

**Trusted entities**

The identity provider(s) ec2.amazonaws.com

**Conditions**

The following conditions define how and when trusted entities can assume the role.

There are no conditions associated with this role.

Click on Edit trust relationship

- Paste this there

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Principal": {
        "Service": "vpc-flow-logs.amazonaws.com"
      },
      "Action": "sts:AssumeRole"
    }
  ]
}
```

6. Go back to your VPC Flow log dashboard and select the role you created. Click on Create flow logs.

## Create flow log



The following flow logs were created:

**Flow Log IDs** fl-07e0cc5a4c1ea4ff1

The output of the flow log created

7. Wait for some time, return back to the CloudWatch dashboard and click on log groups you will see the Log stream Traffic.

```
2 579807160478 eni-07ec21a29debdea6a 172.31.35.55 18.185.97.81 80 40642 6 80 52700 1534192449 1534192626 ACCEPT OK
2 579807160478 eni-07ec21a29debdea6a 18.185.97.81 172.31.35.55 40096 80 6 80 5800 1534192449 1534192626 ACCEPT OK
2 579807160478 eni-07ec21a29debdea6a 18.185.97.81 172.31.35.55 40398 80 6 79 5748 1534192449 1534192626 ACCEPT OK
2 579807160478 eni-07ec21a29debdea6a 172.31.35.55 18.185.97.81 80 41668 6 81 52752 1534192449 1534192626 ACCEPT OK
2 579807160478 eni-07ec21a29debdea6a 18.185.97.81 172.31.35.55 80 40604 6 82 52804 1534192449 1534192626 ACCEPT OK
2 579807160478 eni-07ec21a29debdea6a 172.31.35.55 18.185.97.81 41970 80 6 77 5644 1534192449 1534192626 ACCEPT OK
2 579807160478 eni-07ec21a29debdea6a 18.185.97.81 172.31.35.55 80 41930 6 80 52700 1534192449 1534192626 ACCEPT OK
2 579807160478 eni-07ec21a29debdea6a 172.31.35.55 18.185.97.81 80 41094 6 80 52700 1534192449 1534192626 ACCEPT OK
2 579807160478 eni-07ec21a29debdea6a 172.31.35.55 18.185.97.81 42022 80 6 80 5800 1534192449 1534192626 ACCEPT OK
2 579807160478 eni-07ec21a29debdea6a 172.31.35.55 18.185.97.81 41212 80 6 80 5800 1534192449 1534192626 ACCEPT OK
2 579807160478 eni-07ec21a29debdea6a 172.31.35.55 18.185.97.81 80 40466 6 79 52648 1534192449 1534192626 ACCEPT OK
2 579807160478 eni-07ec21a29debdea6a 172.31.35.55 18.185.97.81 80 41182 6 81 52752 1534192449 1534192626 ACCEPT OK
2 579807160478 eni-07ec21a29debdea6a 172.31.35.55 18.185.97.81 80 39788 6 78 52596 1534192449 1534192626 ACCEPT OK
2 579807160478 eni-07ec21a29debdea6a 172.31.35.55 18.185.97.81 40188 80 6 80 5800 1534192449 1534192626 ACCEPT OK
2 579807160478 eni-07ec21a29debdea6a 172.31.35.55 18.185.97.81 39204 80 6 79 5748 1534192449 1534192626 ACCEPT OK
```

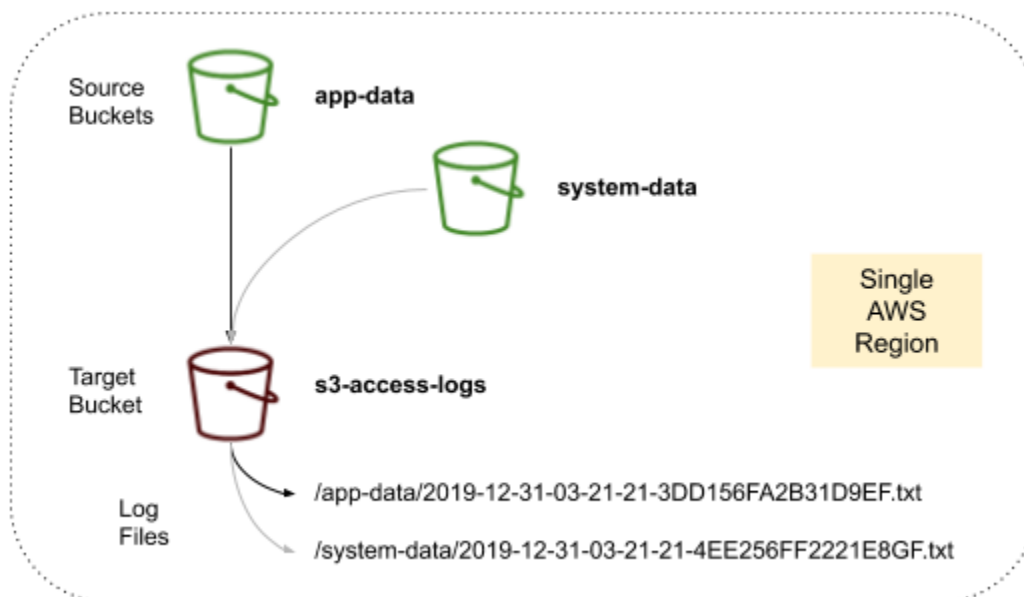
How a typical Flow log looks like

In the Flow log dashboard you can also filter traffic coming from a specific IP address.

Big Ups! You were able to capture information about the IP traffic in your VPC, storing the raw data in Amazon CloudWatch where it can be retrieved and viewed.

[Here](#) is the **official AWS documentation** for the same process we gone through above.

## 2. Setup S3 Access Logs to Log Management S3 Bucket.



### 1. Enabling Logging for Bucket Objects

To use S3 logs, you first need to create one bucket to store files (objects) and another to store the logs. This should be created in the same region. It is a good practice not to save the logs in the same bucket because we want to save the logs for the interactions that the bucket receives and if the bucket has a problem the logs may not be able to be saved with the information about what is causing the error.

Buckets (3) [Info](#)

↻

Copy ARN

Empty

Delete

Create bucket

Find buckets by name

< 1 > ⚙

	Name	AWS Region	Access	Creation date
<input type="radio"/>	mb1987-bucket	US West (N. California) us-west-1	Bucket and objects not public	February 4, 2022, 23:12:20 (UTC-03:00)
<input type="radio"/>	mb1987-bucket-logs	US West (N. California) us-west-1	Bucket and objects not public	February 4, 2022, 23:12:51 (UTC-03:00)

After you've created the buckets, go to the Properties of the bucket that will store the files to associate it with the bucket for logs. On the Properties page, click on the Edit button in the Server access logging box. In this form, select Enable to allow the bucket to provide log data about stored objects, then click on Browse S3 to select the log bucket.

Amazon S3 > mb1987-bucket > Edit server access logging

## Edit server access logging [Info](#)


### Server access logging

Log requests for access to your bucket. [Learn more](#)

Server access logging

☐ Disable

☒ Enable



**Bucket policy will be updated**  
When you enable server access logging, the S3 console automatically updates your bucket policy to include access to the S3 log delivery group.

Target bucket

[Browse S3](#)

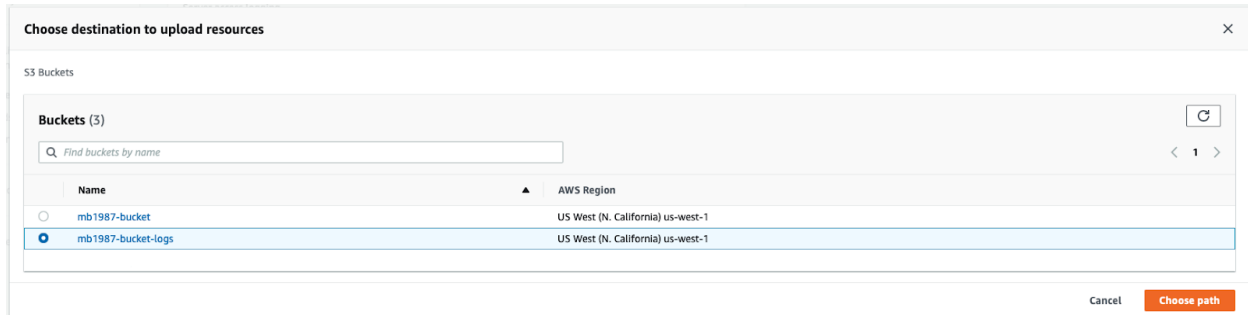
Format: s3://bucket/prefix

Cancel

Save changes

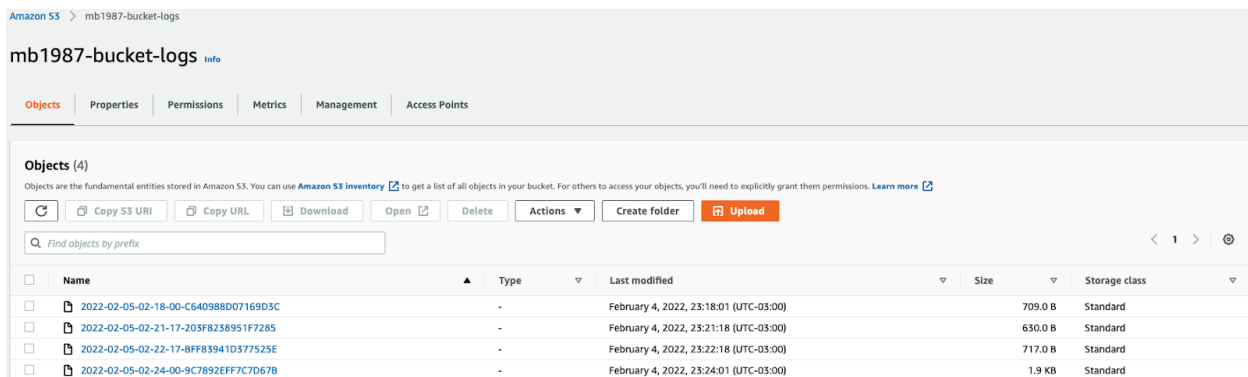


In the modal, select the proper bucket and click on Choose path. Back in the form, click on Save changes to apply the association between the buckets. Clicking that button is all you need to do to start saving object usage logs.



## 2. Testing Logging

Now, let's try to access the bucket and upload some files. For this test, there is no need to add any other settings. You can open, download, and remove files to generate logs for these actions. Then, you can access the bucket for logs and wait a few minutes to receive the log data about the newly uploaded file. Then, open the logs to see the type of data available in the log information. You will see logs related to actions taken within the bucket to get or remove objects, along with policies and versioning information.



## 3. Log Data Samples

Let's take a look at some sample logs and their available formats on AWS S3. For example, this is DELETE:

```
ce6f2c543de2b9a3a4fdf21d56e95135af4045032a9157cb2fdb1a4854c73110 mb1987-bucket [05/Feb/2022:01:09:08 +0000]
191.XXX.XXX.216 ce6f2c543de2b9a3a4fdf21d56e95135af4045032a9157cb2fdb1a4854c73110 BW21GJ60HATK1VCR
REST.POST.MULTI_OBJECT_DELETE - "POST /mb1987-bucket?delete= HTTP/1.1" 200 - 5100 - 802 - "-" "S3Console/0.4,
aws-internal/3 aws-sdk-java/1.11.1030 Linux/5.4.172-100.336.amzn2int.x86_64 OpenJDK_64-Bit_Server_VM/25.302-
b08 java/1.8.0_302 vendor/Oracle_Corporation cfg/retry-mode/standard" -
NLZUuMfnLn6Kq3LKEB5GFRNAVEo9cy/BR5bmzhy4dXWD0ogwa6Q71IzUbJKlidFbVwUiRR9jk9w= SigV4 ECDHE-RSA-AES128-GCM-
SHA256 AuthHeader s3-us-west-1.amazonaws.com TLSv1.2 -
```

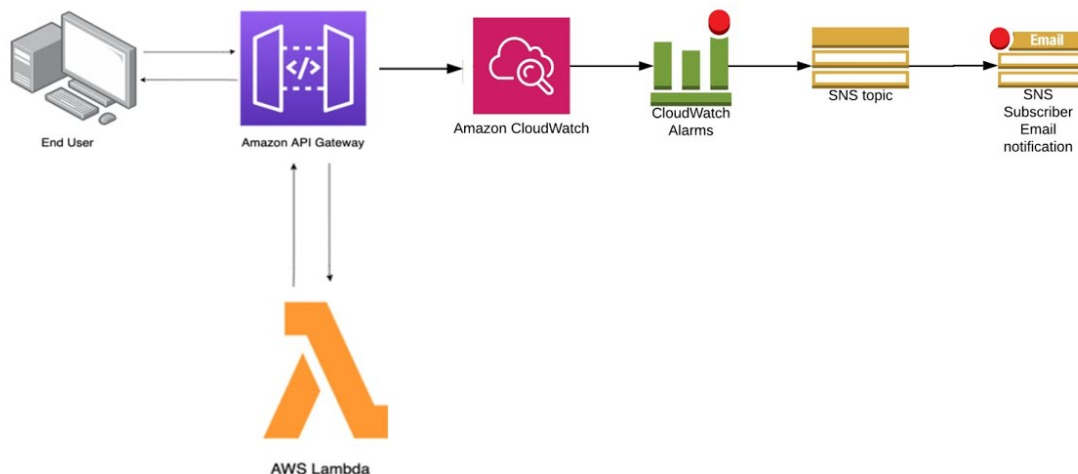
And this is GET:

```
ce6f2c543de2b9a3a4fdf21d56e95135af4045032a9157cb2fdb1a4854c73110 mb1987-bucket [05/Feb/2022:01:08:58 +0000]
191.XXX.XXX.216 ce6f2c543de2b9a3a4fdf21d56e95135af4045032a9157cb2fdb1a4854c73110 MSAHCXNXGA3NJY9V
REST.GET.BUCKET - "GET /mb1987-bucket?list-type=2&encoding-type=url&max-keys=1&fetch-owner=true&delimiter=&prefix= HTTP/1.1" 200 - 768 - 106 105 "-" "S3Console/0.4, aws-internal/3 aws-sdk-java/1.11.1030 Linux/5.4.172-100.336.amzn2int.x86_64 OpenJDK_64-Bit_Server_VM/25.302-b08 java/1.8.0_302 vendor/Oracle_Corporation cfg/retry-mode/standard" -
X4VZjjWLHbLLApCiTeXoulUTQmYYduxisCY8E6se4YZx7FRBX6KsSd4lG+VwfrpgLqQXTdPbTGGE= SigV4 ECDHE-RSA-AES128-GCM-SHA256 AuthHeader s3-us-west-1.amazonaws.com TLSv1.2 -
```

[Here](#) is the **official AWS documentation** for the same process we have gone through above.

### 3. CloudWatch API logging using the API Gateway console

## Monitoring and Logging API Activity



To set up CloudWatch API logging, you must have deployed the API to a stage. You must also have configured [an appropriate CloudWatch Logs role](#) ARN for your account.

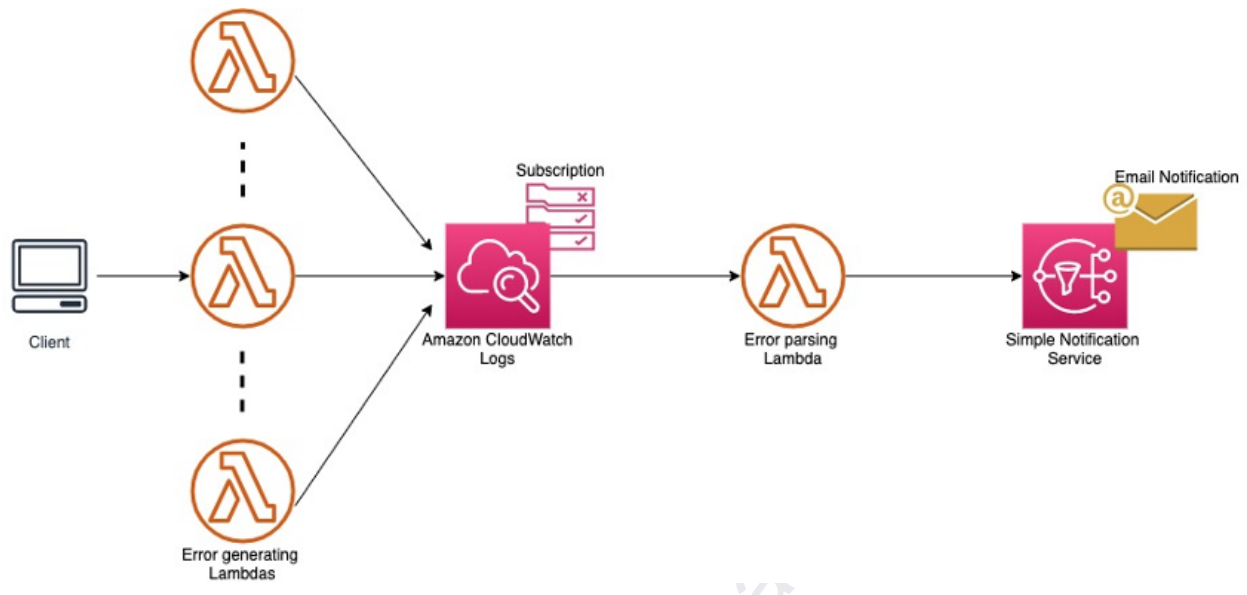
1. Sign in to the API Gateway console at <https://console.aws.amazon.com/apigateway>
2. Choose a REST API.
3. Choose Settings from the primary navigation panel and enter an ARN of an IAM role with appropriate permissions in CloudWatch log role ARN. You need to do this once.

4. Do one of the following:
  - a. Choose an existing API and then choose a stage.
  - b. Create an API and deploy it to a stage.
5. Choose Logs/Tracing in the Stage Editor.
6. To enable execution logging:
  - a. Choose a logging level from the CloudWatch Logs dropdown menu.  
**Warning:** *Full Request and Response Logs can be useful to troubleshoot APIs, but can result in logging sensitive data. We recommend not using Full Request and Response Logs for production APIs.*
  - b. If desired, choose Enable Detailed CloudWatch Metrics.
7. For more information about CloudWatch metrics, see [Monitoring REST API execution with Amazon CloudWatch metrics](#).
8. To enable access logging:
  - a. Choose Enable Access Logging under Custom Access Logging.
  - b. Enter the ARN of a log group in Access Log Destination ARN. The ARN format is `arn:aws:logs:{region}:{account-id}:log-group:log-group-name`.
  - c. Enter a log format in Log Format. You can choose CLF, JSON, XML, or CSV to use one of the provided examples as a guide.
9. Choose to Save Changes.

**Note:** *You can enable execution logging and access logging independently of each other.*

API Gateway is now ready to log requests to your API. You don't need to redeploy the API when you update the stage settings, logs, or stage variables.

## 4. Monitoring AWS Lambda Function Logs in CloudWatch Stream specific to Lambda Function CloudWatch Log Group.



### Prerequisites:

[Execution role](#) attached to the Lambda Function needs permission to upload logs to CloudWatch Logs. You can add CloudWatch Logs permissions using the AWSLambdaBasicExecutionRole AWS managed policy provided by Lambda.

To add this policy to your role, run the following AWS CLI command (*Make sure you have configure AWS CLI using Access Key / Secret Key on the machine where you are executing the below AWS CLI command*):

```
aws iam attach-role-policy --role-name your-role --policy-arn
arn:aws:iam::aws:policy/service-role/AWSLambdaBasicExecutionRole
```

For more information, see [AWS managed policies for Lambda features](#).

### Pricing:

There is no additional charge for using Lambda logs; however, standard CloudWatch Logs charges apply. For more information, see [CloudWatch pricing](#).

### Using the Lambda console to view Lambda execution logs in CloudWatch:

To view logs using the Lambda console

1. Open the [Functions page](#) of the Lambda console.

2. Choose a function.
3. Choose **Monitor**.
4. Choose **View logs in CloudWatch**.


## 5. Capturing EC2 instance CPU/Mem/Disk Utilisation metrics by installing CW Agent on the server and monitoring them using a Custom CW Dashboard.



### Step 1- Launch an EC2 with Ubuntu AMI 20.04

I believe that we can launch its own so I'm not explaining the process for it.

Used AMI: (Ubuntu 20.04)



**Ubuntu Server 20.04 LTS (HVM), SSD Volume Type** - ami-09e67e426f25ce0d7 (64-bit x86) / ami-00d1ab6b335f217cf (64-bit Arm)

Free tier eligible

Ubuntu Server 20.04 LTS (HVM),EBS General Purpose (SSD) Volume Type. Support available from Canonical (<http://www.ubuntu.com/cloud/services>).

Root device type: ebs    Virtualization type: hvm    ENA Enabled: Yes

EC2 Instance: (Type: t2.micro)

Instances (1/1) <a href="#">Info</a>				
<input type="text" value="Filter instances"/>				
Name: Grafana-Integration <span>×</span> <span>Clear filters</span>				
<input checked="" type="checkbox"/>	Name	Instance ID	Instance state	Instance type
<input checked="" type="checkbox"/>	Grafana-Integration	i-0d4e8776dd465e5d7	<span>✓</span> Running <span>🔍</span>	t2.micro

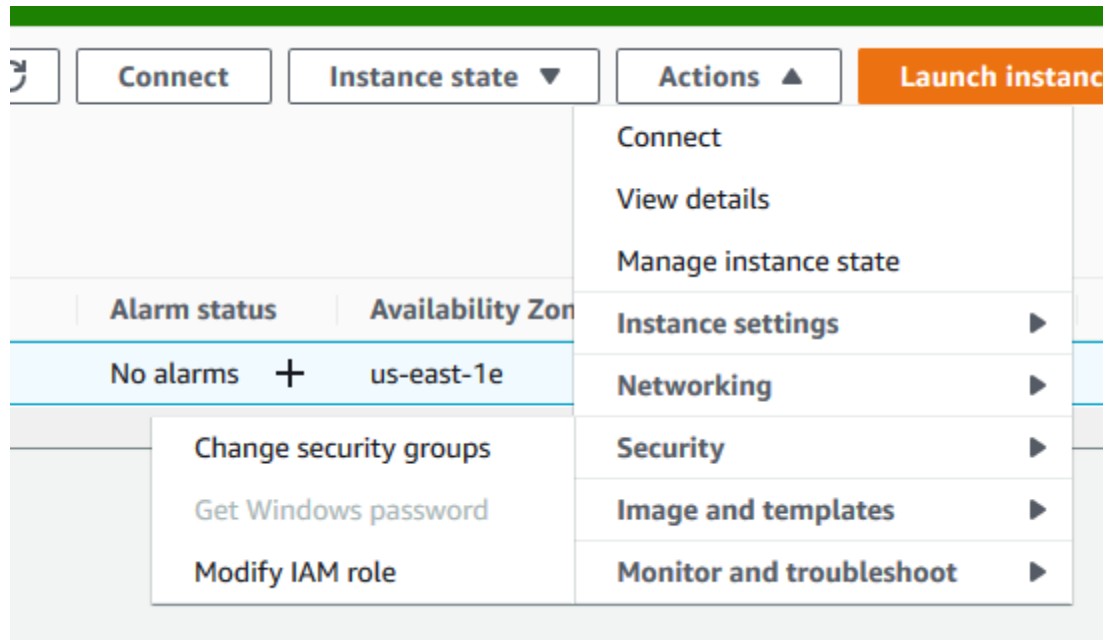
## Step 2: Attach an IAM role with EC2 Instance

Access to AWS resources requires permissions. You create an IAM role, an IAM user, or both to grant permissions that the CloudWatch agent needs to write metrics to CloudWatch. If you're going to use the agent on Amazon EC2 instances, you must create an IAM role. If you're going to use the agent on on-premises servers, you must create an IAM user.

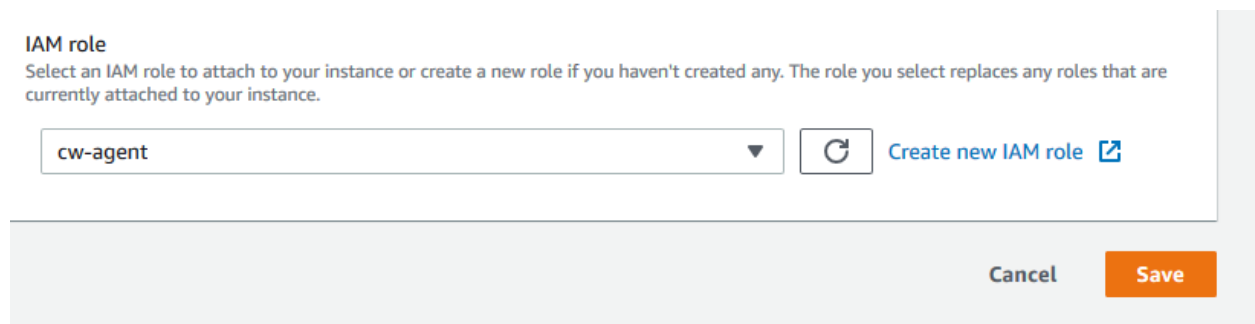
Create an IAM role, Make sure that AWS service is selected under Select type of trusted entity. For Choose a use case, choose EC2 under Common use cases, and Choose Next: Permissions. In the list of policies, select CloudWatchAgentServerPolicy & create an IAM role.

Permissions	Trust relationships	Tags	Access Advisor	Revoke sessions
Permissions policies (1 policy applied)				
<span>Attach policies</span>				
Policy name				Policy type
<span>▶</span> CloudWatchAgentServerPolicy				AWS managed policy

Now attach an IAM role with EC2 Instance. Go to EC2 — Select Instance — Click on Action — Security — Modify IAM role



Choose Create IAM role & save it.



### Step 3: Download the CloudWatch Agent Package

Use the following steps to download the CloudWatch agent package, SSH to Instance & Download CW Agent package.

Download the CloudWatch agent:

```
wget
```

```
https://s3.amazonaws.com/amazoncloudwatch-agent/ubuntu/amd64/latest/amazon-cloudwatch-agent.deb
```

```

root@ip-172-31-52-40:/home/ubuntu# wget https://s3.amazonaws.com/amazoncloudwatch-agent/ubuntu/amd64/latest/amazon-cloudwatch-agent.deb
--2021-05-04 05:34:37-- https://s3.amazonaws.com/amazoncloudwatch-agent/ubuntu/amd64/latest/amazon-cloudwatch-agent.deb
Resolving s3.amazonaws.com (s3.amazonaws.com)... 52.216.1.27
Connecting to s3.amazonaws.com (s3.amazonaws.com)|52.216.1.27|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 54662644 (52M) [application/octet-stream]
Saving to: 'amazon-cloudwatch-agent.deb'

amazon-cloudwatch-agent.deb      100%[=====] 52.13M  47.0MB/s   in 1.1s

2021-05-04 05:34:38 (47.0 MB/s) - 'amazon-cloudwatch-agent.deb' saved [54662644/54662644]
root@ip-172-31-52-40:/home/ubuntu#

```

Install the package:

`dpkg -i -E ./amazon-cloudwatch-agent.deb`

```

root@ip-172-31-52-40:/home/ubuntu# dpkg -i -E ./amazon-cloudwatch-agent.deb
Selecting previously unselected package amazon-cloudwatch-agent.
(Reading database ... 60149 files and directories currently installed.)
Preparing to unpack ./amazon-cloudwatch-agent.deb ...
create group cwagent, result: 0
create user cwagent, result: 0
Unpacking amazon-cloudwatch-agent (1.247347.6b250880-1) ...
Setting up amazon-cloudwatch-agent (1.247347.6b250880-1) ...
root@ip-172-31-52-40:/home/ubuntu#

```

Update Packages & Install collectd: (This will take a few minutes if you haven't updated your available updates prior)

`apt-get update && apt-get install collectd`

```

0 upgraded, 557 newly installed, 0 to remove and 4 not upgraded.
Need to get 216 MB of archives.
After this operation, 1150 MB of additional disk space will be used.
Do you want to continue? [Y/n] Y
Get:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu focal-updates/main amd64 libsane-common all 1.0.29-0ubuntu5.2 [277 kB]
Get:2 http://us-east-1.ec2.archive.ubuntu.com/ubuntu focal/main amd64 libtalloc2 amd64 2.3.0-3ubuntu1 [29.5 kB]
Get:3 http://us-east-1.ec2.archive.ubuntu.com/ubuntu focal/main amd64 libevent0 amd64 0.10.1-4 [35.5 kB]
Get:4 http://us-east-1.ec2.archive.ubuntu.com/ubuntu focal-updates/main amd64 libwbclient0 amd64 2:4.11.6+dfsg-0ubuntu1.8 [221 kB]

```

```

aspell-autobuildhash: processing: en [en_US-w_accents-only].
aspell-autobuildhash: processing: en [en_US-wo_accents-only].
Processing triggers for libgdk-pixbuf2.0-0:amd64 (2.40.0+dfsg-3ubuntu0.2) ...
Processing triggers for rygel (0.38.3-1ubuntu1) ...
Processing triggers for libc-bin (2.31-0ubuntu9.2) ...
Processing triggers for dbus (1.12.16-2ubuntu2.1) ...
Processing triggers for systemd (245.4-4ubuntu3.6) ...
Processing triggers for sgml-base (1.29.1) ...
root@ip-172-31-52-40:/home/ubuntu#

```

## Step 4: Create the CloudWatch Agent Configuration File

Before running the CloudWatch agent on any servers, you must create a CloudWatch agent configuration file. The agent configuration file is a JSON file that specifies the



metrics and logs that the agent is to collect, including custom metrics. The agent configuration file wizard, amazon-cloud watch-agent-config-wizard, asks for a series of questions.

/opt/aws/amazon-cloudwatch-agent/bin/amazon-cloudwatch-agent-config-wizard

We're using Linux machine so option 1

```
root@ip-172-31-52-40:/home/ubuntu# /opt/aws/amazon-cloudwatch-agent/bin/amazon-cloudwatch-agent-config-wizard
=====
= Welcome to the AWS CloudWatch Agent Configuration Manager =
=====
On which OS are you planning to use the agent?
1. linux
2. windows
3. darwin
default choice: [1]:
```

Using EC2 Instance so Option 1

```
Trying to fetch the default region based on ec2 metadata...
Are you using EC2 or On-Premises hosts?
1. EC2
2. On-Premises
default choice: [1]:
```

Select a user

```
Which user are you planning to run the agent?
1. root
2. cwagent
3. others
default choice: [1]:
```

We can skip these two options for memory metric.

```

Do you want to turn on StatsD daemon?
1. yes
2. no
default choice: [1]:
2
Do you want to monitor metrics from CollectD?
1. yes
2. no
default choice: [1]:
2

```

We can choose the below options according to our requirements.

```

Do you want to monitor any host metrics? e.g. CPU, memory, etc.
1. yes
2. no
default choice: [1]:
1
Do you want to monitor cpu metrics per core? Additional CloudWatch charges may apply.
1. yes
2. no
default choice: [1]:
2
Do you want to add ec2 dimensions (ImageId, InstanceId, InstanceType, AutoScalingGroupName) into all of your metrics if the info is available?
1. yes
2. no
default choice: [1]:
2
Would you like to collect your metrics at high resolution (sub-minute resolution)? This enables sub-minute resolution for all metrics, but you
cific metrics in the output json file.
1. 1s
2. 10s
3. 30s
4. 60s
default choice: [4]:
4
Which default metrics config do you want?
1. Basic
2. Standard
3. Advanced
4. None
default choice: [1]:
1

```

Provide some declarations for config files.

```

Are you satisfied with the above config? Note: it can be manually customized after the wizard completes to add additional items.
1. yes
2. no
default choice: [1]:
1
Do you have any existing CloudWatch Log Agent (http://docs.aws.amazon.com/AmazonCloudWatch/latest/logs/AgentReference.html) confi
on?
1. yes
2. no
default choice: [2]:
2
Do you want to monitor any log files?
1. yes
2. no
default choice: [1]:
2
Saved config file to /opt/aws/amazon-cloudwatch-agent/bin/config.json successfully.
Current config as follows:

```

We can check json file under /opt/aws/amazon-cloudwatch-agent/bin/config.json

```
Saved config file to /opt/aws/amazon-cloudwatch-agent/bin/config.json successfully.
Current config as follows:
{
  "agent": {
    "metrics_collection_interval": 60,
    "run_as_user": "root"
  },
  "metrics": {
    "metrics_collected": {
      "disk": {
        "measurement": [
          "used_percent"
        ],
        "metrics_collection_interval": 60,
        "resources": [
          "*"
        ]
      },
      "mem": {
        "measurement": [
          "mem_used_percent"
        ],
        "metrics_collection_interval": 60
      }
    }
  }
}
Please check the above content of the config.
The config file is also located at /opt/aws/amazon-cloudwatch-agent/bin/config.json.
Edit it manually if needed.
Do you want to store the config in the SSM parameter store?
1. yes
2. no
default choice: [1]:
2
Program exits now.
```

Now check the status of CW agent.

/opt/aws/amazon-cloudwatch-agent/bin/amazon-cloudwatch-agent-ctl -m ec2 -a status

OR

systemctl status amazon-cloudwatch-agent

```
root@ip-172-31-52-40:/home/ubuntu# /opt/aws/amazon-cloudwatch-agent/bin/amazon-cloudwatch-agent-ctl -m ec2 -a status
{
  "status": "stopped",
  "starttime": "",
  "configstatus": "not configured",
  "cwoc_status": "stopped",
  "cwoc_starttime": "",
  "cwoc_configstatus": "not configured",
  "version": "1.247347.6b250880"
}
root@ip-172-31-52-40:/home/ubuntu#
```

Status is stopped now, so start it & check status after that.

```
/opt/aws/amazon-cloudwatch-agent/bin/amazon-cloudwatch-agent-ctl -a fetch-config -m ec2 -s -c file:/opt/aws/amazon-cloudwatch-agent/bin/config.json
```

```
root@ip-172-31-52-40:/home/ubuntu# /opt/aws/amazon-cloudwatch-agent/bin/amazon-cloudwatch-agent-ctl -a fetch-config -m ec2 -s -c file:/opt/aws/amazon-cloudwatch-agent/bin/config.json
***** processing amazon-cloudwatch-agent *****
/opt/aws/amazon-cloudwatch-agent/bin/config-downloader --output-dir /opt/aws/amazon-cloudwatch-agent/etc/amazon-cloudwatch-agent.d --download-source file:/opt/aws/amazon-cloudwatch-agent/bin/config.json --mode ec2 --config /opt/aws/amazon-cloudwatch-agent/etc/common-config.toml --multi-config default
Successfully fetched the config and saved in /opt/aws/amazon-cloudwatch-agent/etc/amazon-cloudwatch-agent.d/file_config.json.tmp
Start configuration validation...
/opt/aws/amazon-cloudwatch-agent/bin/config-translator --input /opt/aws/amazon-cloudwatch-agent/etc/amazon-cloudwatch-agent.json --input-dir /opt/aws/amazon-cloudwatch-agent/etc/amazon-cloudwatch-agent.d --output /opt/aws/amazon-cloudwatch-agent/etc/amazon-cloudwatch-agent.toml --mode ec2 --config /opt/aws/amazon-cloudwatch-agent/etc/common-config.toml --multi-config default
2021/05/04 05:59:30 Reading json config file path: /opt/aws/amazon-cloudwatch-agent/etc/amazon-cloudwatch-agent.d/file_config.json.tmp ...
Valid json input schema.
!! Detecting run as user...
No csm configuration found.
No log configuration found.
Configuration validation first phase succeeded
/opt/aws/amazon-cloudwatch-agent/bin/amazon-cloudwatch-agent -schematest -config /opt/aws/amazon-cloudwatch-agent/etc/amazon-cloudwatch-agent.toml
Configuration validation second phase succeeded
Configuration validation succeeded
amazon-cloudwatch-agent has already been stopped
Created symlink /etc/systemd/system/multi-user.target.wants/amazon-cloudwatch-agent.service → /etc/systemd/system/amazon-cloudwatch-agent.service.
root@ip-172-31-52-40:/home/ubuntu# /opt/aws/amazon-cloudwatch-agent/bin/amazon-cloudwatch-agent-ctl -m ec2 -a status
{
  "status": "running",
  "starttime": "2021-05-04T05:59:31+00:00",
  "configstatus": "configured",
  "cwoc_status": "stopped",
  "cwoc_starttime": "",
  "cwoc_configstatus": "not configured",
  "version": "1.247347.6b250880"
}
root@ip-172-31-52-40:/home/ubuntu#
```

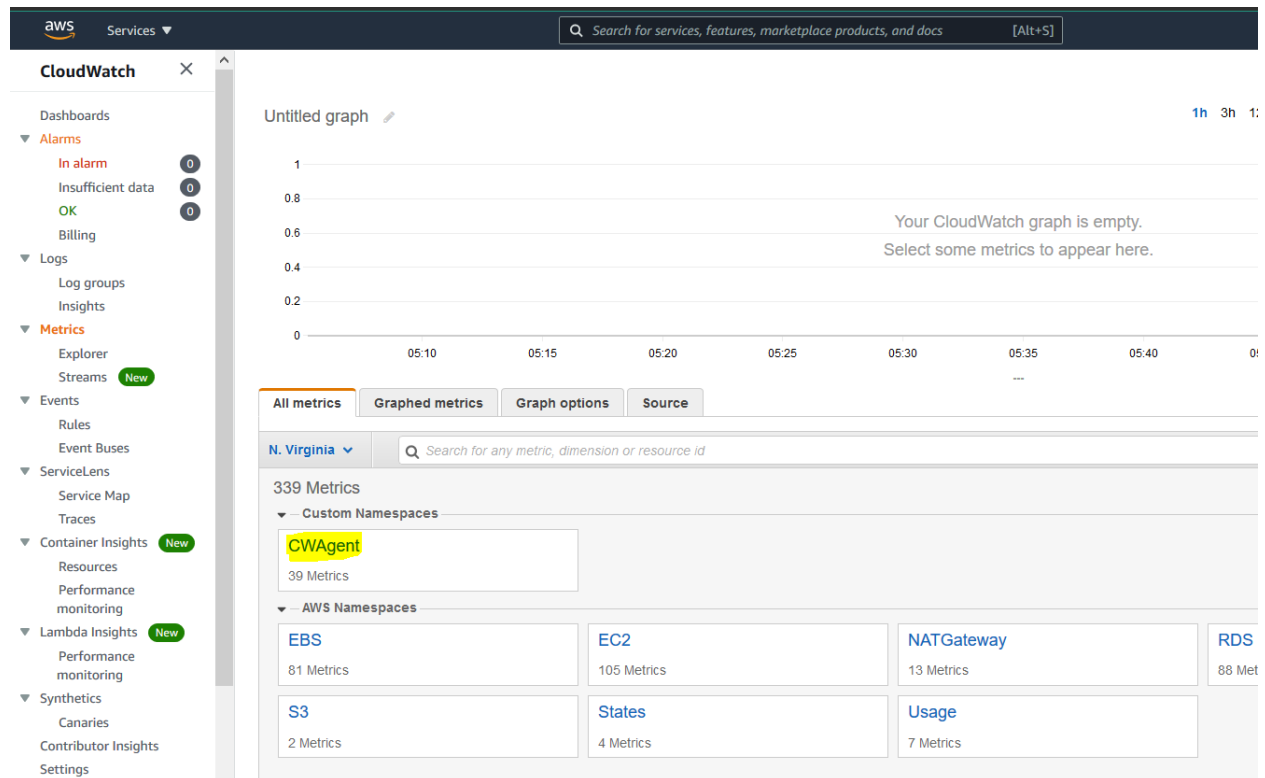
## Troubleshooting:

If config file for CW Agents emits error, execute below commands to setup collectd service for CW Agent.

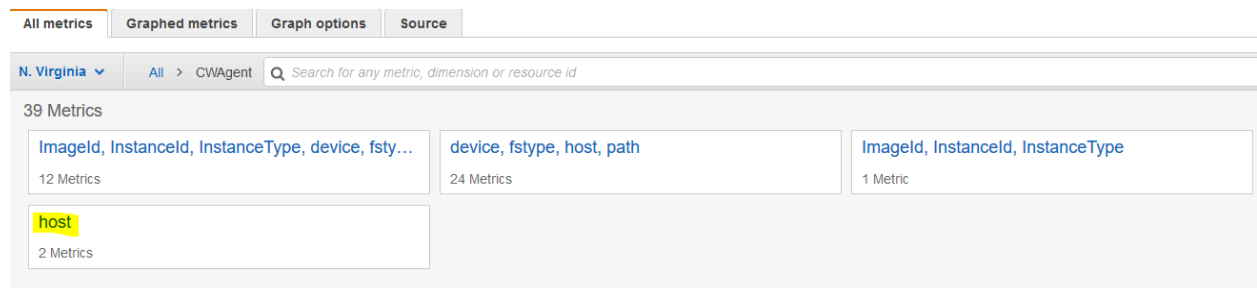
```
sudo apt install collectd -y
mkdir -p /usr/share/collectd/
touch /usr/share/collectd/types.db
```

## Step 5: Check Custom Metrics in AWS CloudWatch:

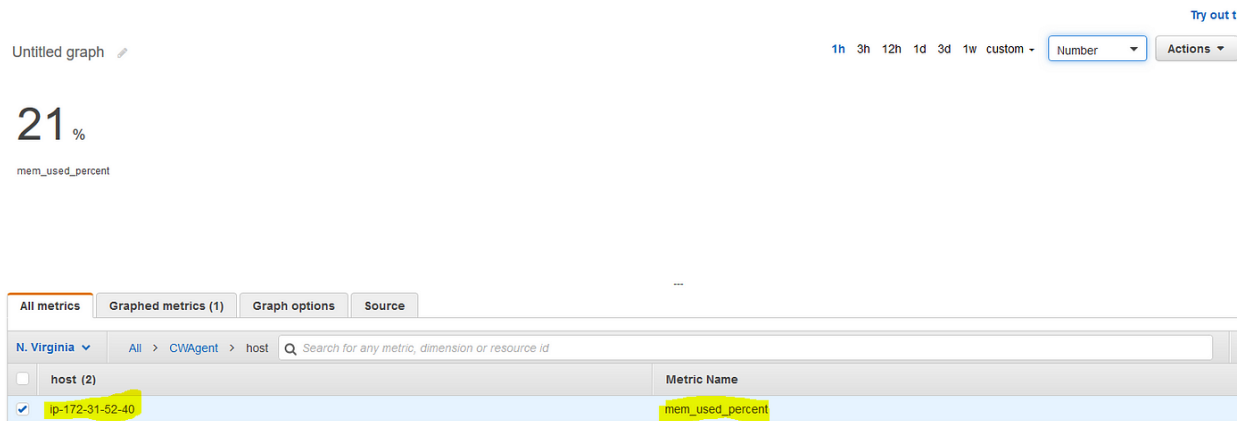
### Custom Metrics:



### Check inside a host for memory metrics:



Select the host according to the private IP of the instance & here you go with the memory metrics of the instance.



## Step 6: Create a custom CW Dashboard and plot metrics on the created dashboard to monitor:

Now, since the metrics are visible on the CW Metrics, you can create a Custom Dashboard with Line graph, Bar graph, Pie chart visualization as per your monitoring requirements.

Select the metrics of your choice (which you want to monitor on the custom dashboard) under CWAgent Custom namespaces and click on Actions → Add to Dashboard → Select your custom dashboard where you want to capture and monitor these metrics.


## 6. Track user activities on your Cloud Environment using CloudTrail on any AWS Cloud Service.

For example, let's assume a scenario: You are given the task to monitor all the user activities performed by the Cloud, DevOps, and Dev team on the AWS S3 bucket service. You can leverage CloudTrail to track API calls made via Console, CLI, SDK, or any other means.

So, all you need to do is just enter CloudTrail in the AWS service search box and go to the CloudTrail service. Click on **Create Trail** and then provide the basic details like Trail name, Storage bucket name, and folder, whether you want to log the user activities in CW Log Groups as well or not, and Tag values.

Once you create the Trail, go to the newly created Trail from the home page of CloudTrail and navigate to the **Data Events** section. Click on the Edit button and you can select as per the requirements such as Event type, Log Selector template (to select logging level), etc., and create a new Data event to capture in the Target Trail S3 bucket.

## Data events [Info](#)

Data events show information about the resource operations performed on or within a resource. [Additional charges apply](#) 



### Advanced event selectors are enabled

Use the following fields for fine-grained control over the data events captured by your trail.

[Switch to basic event selectors](#)

#### ▼ Data event: S3

[Remove](#)

#### Data event type

Choose the source of data events to log.

S3

#### Log selector template

Log all events

#### Selector name - *optional*

*Enter a name*

1,000 character limit

#### ► JSON view

[Add data event type](#)

[Cancel](#)

[Save changes](#)

Finally, after applying these configurations, you can start receiving logs in the target S3 bucket or/and the AWS CW Log Group as per your selection within a few minutes as soon as the configured data event is triggered (In this case we selected S3 All events to capture) in a JSON format as below:

Prep

```
    },
    {
      "eventVersion": "1.08",
      "userIdentity": {
        "type": "Root",
        "principalId": "615311846444",
        "arn": "arn:aws:iam::615311846444:root",
        "accountId": "615311846444",
        "accessKeyId": "ASIAY6Q3RVQW0JTZLX3J",
        "sessionContext": {
          "sessionIssuer": {},
          "webIdFederationData": {},
          "attributes": {
            "creationDate": "2023-04-28T00:22:24Z",
            "mfaAuthenticated": "true"
          }
        }
      },
      "eventTime": "2023-04-28T01:27:51Z",
      "eventSource": "s3.amazonaws.com",
      "eventName": "GetBucketAcl",
      "awsRegion": "ca-central-1",
      "sourceIPAddress": "99.231.111.68",
      "userAgent": "[S3Console/0.4, aws-internal/3 aws-sdk-java/1.11.1030 Linux/5.4.238-155.347.amzn2int.x86_64 OpenJDK_64-Bit_Server_VM/25.362-b10 java/1.8.0_362 vendor/Oracle_Corporation cfg/retry-mode/standard]",
      "requestParameters": {
        "bucketName": "test-bucket-2456456",
        "Host": "test-bucket-2456456.s3.ca-central-1.amazonaws.com",
        "acl": ""
      }
    }
  ],
  "responseElements": {
    "acl": ""
  }
}
```

Prepared by Dhru